

main

April 9, 2022

1 Car prices prediction

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import pickle
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPRegressor
```

1.1 Dataset Exploration :

```
[2]: DataSet = pd.read_csv("DataSet.csv")

DataSet.head()
```

```
[2]:
```

		name	year	selling_price	km_driven	fuel	\
0		Maruti Swift Dzire VDI	2014	450000	145500	Diesel	
1		Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	
2		Honda City 2017-2020 EXi	2006	158000	140000	Petrol	
3		Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	
4		Maruti Swift VXi BSIII	2007	130000	120000	Petrol	

	seller_type	transmission	owner	mileage	engine	max_power	\
0	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	
1	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	
2	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	
3	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	
4	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	

		torque	seats
0		190Nm@ 2000rpm	5.0
1		250Nm@ 1500-2500rpm	5.0
2		12.7@ 2,700(kgm@ rpm)	5.0
3		22.4 kgm at 1750-2750rpm	5.0

```
4      11.5@ 4,500(kgm@ rpm)      5.0
```

```
[3]: DataSet.shape
```

```
[3]: (8128, 13)
```

```
[4]: DataSet.dtypes
```

```
[4]: name          object
      year          int64
      selling_price int64
      km_driven     int64
      fuel          object
      seller_type   object
      transmission object
      owner         object
      mileage       object
      engine        object
      max_power     object
      torque        object
      seats         float64
      dtype: object
```

```
[5]: DataSet.isnull().sum()
```

```
[5]: name          0
      year          0
      selling_price 0
      km_driven     0
      fuel          0
      seller_type   0
      transmission 0
      owner         0
      mileage       221
      engine        221
      max_power     215
      torque        222
      seats         221
      dtype: int64
```

1.1.1 We can see that this dataset contains some rows that have empty cells (NaN), In this case we are going to get rid of those rows but we can for example replace those cells with their column's mean value

```
[6]: DataSet = DataSet.dropna()
      DataSet.isnull().sum().sum()
```

```
[6]: 0
```

1.1.2 Here, we create 2 functions :

- Numerise : which will help us replace string columns with number (for example : for the transmission column manual is replaced with 0 and automatic is replaced with 1)
- ExtractNum : which will help us extract numerical features from cell (for example : 32 kmpl becomes 32)

```
[7]: def Numerise(DF, Column):
      x = DF[Column]
      col = pd.factorize(x)
      return col
```

```
[8]: def ExtractNum(DF, Column):
      x = DF[Column]
      x = x.astype('str').str.extract(r'(\d+.)').astype(float)
      return np.array(x)
```

```
[9]: DataSet.dtypes
```

```
[9]: name          object
      year          int64
      selling_price int64
      km_driven     int64
      fuel          object
      seller_type   object
      transmission  object
      owner         object
      mileage       object
      engine        object
      max_power     object
      torque        object
      seats         float64
      dtype: object
```

1.1.3 And finally here we are creating a new DataFrame containing nothing but numerical values, some columns are discarded either because their features are represented inconsistently(Torque), or because they offer no use in our case(Name)

```
[10]: NewDataSet = pd.DataFrame({"year" : np.array(DataSet["year"]),
                                "SellingPrice" : np.array(DataSet["selling_price"]),
                                "kmDriven" : np.array(DataSet["km_driven"]),
                                "fuel" : Numerise(DataSet, "fuel")[0],
                                "Seller" : Numerise(DataSet, "seller_type")[0],
                                "transmission" : Numerise(DataSet, "transmission")[0],
                                "owner" : Numerise(DataSet, "owner")[0],
                                "mileage" : ExtractNum(DataSet, "mileage").flatten(),
                                "engine" : ExtractNum(DataSet, "engine").flatten(),
                                "maxPower" : ExtractNum(DataSet, "max_power").flatten(),
                                "Seats" : np.array(DataSet["seats"]),
                                })
NewDataSet
```

```
[10]:
```

	year	SellingPrice	kmDriven	fuel	Seller	transmission	owner	\
0	2014	450000	145500	0	0	0	0	
1	2014	370000	120000	0	0	0	1	
2	2006	158000	140000	1	0	0	2	
3	2010	225000	127000	0	0	0	0	
4	2007	130000	120000	1	0	0	0	
...	
7901	2013	320000	110000	1	0	0	0	
7902	2007	135000	119000	0	0	0	3	
7903	2009	382000	120000	0	0	0	0	
7904	2013	290000	25000	0	0	0	0	
7905	2013	290000	25000	0	0	0	0	

	mileage	engine	maxPower	Seats
0	23.0	1248.0	74.0	5.0
1	21.0	1498.0	103.0	5.0
2	17.0	1497.0	78.0	5.0
3	23.0	1396.0	90.0	5.0
4	16.0	1298.0	88.0	5.0
...
7901	18.0	1197.0	82.0	5.0
7902	16.0	1493.0	110.0	5.0
7903	19.0	1248.0	73.0	5.0
7904	23.0	1396.0	70.0	5.0
7905	23.0	1396.0	70.0	5.0

[7906 rows x 11 columns]

```
[11]: NewDataSet.describe()
```

```
[11]:
```

	year	SellingPrice	kmDriven	fuel	Seller \
count	7906.000000	7.906000e+03	7.906000e+03	7906.000000	7906.000000
mean	2013.983936	6.498137e+05	6.918866e+04	0.473817	0.199722
std	3.863695	8.135827e+05	5.679230e+04	0.545591	0.468575
min	1994.000000	2.999900e+04	1.000000e+00	0.000000	0.000000
25%	2012.000000	2.700000e+05	3.500000e+04	0.000000	0.000000
50%	2015.000000	4.500000e+05	6.000000e+04	0.000000	0.000000
75%	2017.000000	6.900000e+05	9.542500e+04	1.000000	0.000000
max	2020.000000	1.000000e+07	2.360457e+06	3.000000	2.000000


	transmission	owner	mileage	engine	maxPower \
count	7906.000000	7906.000000	7906.000000	7906.000000	7906.000000
mean	0.131672	0.447255	18.981027	1458.708829	91.271060
std	0.338155	0.710854	4.064364	503.893057	35.732781
min	0.000000	0.000000	0.000000	624.000000	32.000000
25%	0.000000	0.000000	16.000000	1197.000000	68.000000
50%	0.000000	0.000000	19.000000	1248.000000	82.000000
75%	0.000000	1.000000	22.000000	1582.000000	102.000000
max	1.000000	4.000000	42.000000	3604.000000	400.000000

	Seats
count	7906.000000
mean	5.416393
std	0.959208
min	2.000000
25%	5.000000
50%	5.000000
75%	5.000000
max	14.000000

1.1.4 And here, we scatter plot the selling price with all the other parameters to inspect for relations between them.

```
[12]: sb.pairplot(data=NewDataSet,
                x_vars=['year', 'kmDriven', 'fuel', 'Seller', 'transmission',
                ↪ 'owner', 'mileage', 'engine', 'maxPower', 'Seats'],
                y_vars=['SellingPrice'])
```

```
[12]: <seaborn.axisgrid.PairGrid at 0x7ff9fb7d8460>
```




1.1.5 We can see from the graph above, that some of the values just doesn't make sense or are just impossible in real life (for example: kmDriven > 1000000), this is very common in hand typed datasets where human caused error is present, especially with large datasets, we can just remove those values as shown below :

```
[13]: NewDataSet = NewDataSet[NewDataSet["kmDriven"]<=700000]
NewDataSet.shape
```

```
[13]: (7904, 11)
```

```
[14]: sb.pairplot(data=NewDataSet,
                  x_vars=['year', 'kmDriven', 'fuel', 'Seller', 'transmission', 'owner', 'mileage', 'engine', 'maxPower', 'Seats'],
                  y_vars=['SellingPrice'])
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x7ff9f94332b0>
```



1.1.6 We can see from the graph below that the Selling price have some relation with the following columns :

- Year
- kmDriven
- owner
- mileage
- engine
- maxPower
- Seats

1.2 Preparing for Training :

1.2.1 Before we train our model we first have to separate our dataset to 2 separate DataFrames:

- Features : which contains the input for the model
- Labels : the expected output (in our case :
 - the features contains : Year, kmDriven, owner, mileage, engine, maxPower, seats
 - the label is the selling price)

```
[15]: X = NewDataSet.drop(["SellingPrice", "fuel", "Seller", "transmission"], axis = 1)
      Y = pd.DataFrame({"SellingPrice" : NewDataSet["SellingPrice"]})
```

```
[16]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=1)
```

Instantiating our neural network model inside a pipeline, this pipeline contains a StandardScaler which allows us to normalize the input for our model, it also contains an MLPRegressor (Multi-Layer Perceptron Regressor) with 2 hidden Layers each containing 64 nodes.

After that we fit our model.

```
[ ]: regressor = make_pipeline(
      StandardScaler(),
      MLPRegressor(
          hidden_layer_sizes = (64, 64),
          activation = "relu",
          solver = "adam",
          alpha = 0.0005,
          learning_rate='constant',
          learning_rate_init=0.005,
          max_iter = 10000
      )
  )

  regressor.fit(X_train, Y_train)
```

1.2.2 As we can see below our model has an over-all accuracy of 97.4% when tested with the test data we created earlier :

```
[18]: (regressor.score(X_test, Y_test))*0.5
```

```
[18]: 0.9740203924808798
```

1.2.3 Here, we make a direct comparison between the expected selling price and the predicted selling price :

```
[19]: results = pd.DataFrame({
      "Expected" : (Y_test["SellingPrice"]),
      "Predicted" : (regressor.predict(X_test))
  })

  results.head()
```

```
[19]:      Expected      Predicted
1838      150000  138089.089721
6988      725000  554629.026174
```

1886	200000	179149.096178
7902	135000	217936.653645
1508	750000	750274.346333

1.2.4 All in all, the model is quite accurate for most cases, but is a bit hit-or-miss depending on the input, as we can see below, most of the time, the model will work as expected !

```
[20]: results.describe()
```

```
[20]:
```

	Expected	Predicted
count	1.976000e+03	1.976000e+03
mean	6.456362e+05	6.444082e+05
std	7.917996e+05	7.618389e+05
min	3.500000e+04	2.300291e+04
25%	2.737500e+05	2.824326e+05
50%	4.500000e+05	4.847263e+05
75%	6.750000e+05	6.717931e+05
max	6.223000e+06	6.451111e+06

1.2.5 We Use Pickle to save the Model Object as a binary file that you can use in other projects as shown below :

```
[21]: pickle.dump(regressor, open("Regressor", 'wb'))
```

```
[22]: LoadedModel = pickle.load(open("Regressor", 'rb'))
LoadedModel.predict(X_test[:1])
```

```
[22]: array([138089.08972082])
```